

# Ch. 1 Overview of Core Audio

---

- engine behind sound on mac or iOS
  - C API
  - exposed in Objective-C, C++, and Swift
- high-level when compared to general audio programming environments
- low-level from an application developer standpoint
- very different from Cocoa and other C-based APIs (Quartz, Core Foundation)
- when not to use
  - simple file playback
  - use AppKit and NSSound or AVAudioPlayer from AV Foundation
  - can also use QTKit on mac

## The Core Audio Frameworks

---

- two groups
  - audio engines - audio stream processing
  - helper APIs - facilitate audio IO to and from these frameworks
- mac and iOS audio engine APIs
  - Audio Units
    - Core Audio passes audio buffers to graphs of Audio Units which do work on the buffers
  - Audio Queues
    - abstraction on top of audio units that simplifies playback and recording
    - uses callbacks
  - OpenAL
    - industry standard 3D audio API (designed to resemble OpenGL)
- helper APIs
  - Audio File Services
    - abstracts access to various file types
  - Audio File Stream Services
    - network and other stream sources of audio
  - Audio Converter Services
    - format conversion
  - Extended Audio File Services
    - allows simultaneous reading and writing and conversion
  - Core MIDI
    - MIDI access
  - Audio Session Services (iOS-only)
    - coordinates use of resources with system on iOS

# Core Audio Conventions

---

- C function calls
  - brings a number of difficulties that modern programmers may not be equipped to deal with
- does mimic object orientation within C in many ways
- Core Foundation is the main C framework and underlies Foundation
  - many of the Objective-C components of Core Foundation gain their functionality by calling C
- Core Audio is very similar in design to Core Foundation

## Your First Core Audio Application

---

- uses Objective-C in Xcode
- see code for details
- notes
  - need to convert from C strings to NSString/CFStringRef
    - can specify UTF-8
    - can use stringByExpandingTildeInPath
  - must convert file paths to NSURL to use in audio file APIs
  - use local variable reference of type AudioFileID type to refer to audio file objects
  - use OSStatus to signal failure or success through return types
  - should check on all Core Audio calls
  - noErr (0) signals success
  - calls AudioFileOpenURL to open file
  - in Objective-C, NSURL can be toll-free cast to CFURLRef
  - to get metadata, ask for kAudioFilePropertyInfoDictionary
    - requires advance allocation of memory for returned metadata
  - to get size for allocation, call AudioFileGetPropertyInfo with AudioFileID
  - refer to docs and headers for info on the various results of getting properties
  - call AudioFileGetProperty
  - Core Foundation does not offer autorelease so retain/release counts must be managed manually
  - AudioFileID has its own clean-up call, AudioFileClose
  - many Core Audio objects have custom clean-up calls, be aware of these
  - check return codes on Core Audio clean-up calls

## Running the Example

---

- requires linking in frameworks in Xcode
- requires passing command-line arguments

- NOTE: the metadata results will vary wildly depending on files/file types/systems
- NOTE: Core Audio requires an adherence to a large number of implicit contracts between various types of calls
  - BE PATIENT AND WORK WITH THE CODE

## Core Audio Properties

---

- preparing and calling property getters and setters in Core Audio is an essential process
- properties
  - key-value pairs
  - keys are enumerated integers
  - values are types defined by APIs
  - values that are retrieved or set are dependent on the specific property
- for AudioFileGetProperty, the keys are 32-bit integers that can be represented as four character codes
  - NOTE: use single-quotes to represent char literals in C
- many Core Audio types have related \*GetProperty() functions
- function signature adds information based on naming conventions of function parameters
  - in - only used for input to the function
  - out - only used for output from the function
  - io - used for input and output from a function

## Ch. 2 The Story of Sound

---

### Making Waves

---

### Digital Audio

---

### DIY Samples

---

- need to set
  - sample rate
  - length of audio sample to create
  - filename/path of output file
  - file format
  - frequency of audio signal (as command-line arg)
  - generate full filename from file format
- need to populate info in AudioStreamBasicDescription struct
  - channels
  - format

- bitrate
- etc
- initialize remaining fields of `AudioStreamBasicDescription` to 0 (using `memset` in C)
  - Core Audio populates certain fields depending on format
- set info in `AudioStreamBasicDescription`
  - single channel
  - PCM
  - 44100
  - 16-bit
  - 2 bytes per frame (one channel \* 2 bytes of sample data)
  - endianness (for PCM)
  - numeric format of samples (`kAudioFormatFlagsSignedInteger`)
  - packed (sample uses all bits available in each byte) vs unpacked
- packets are only useful for VBR formats
  - `bytesPerFrame` and `bytesPerPacket` are equal for CBR formats
- AIFF only handles big endian PCM
- numeric format and packed vs unpacked are flags that must be bitwise or'ed together
- create `AudioFileID` for audio file using `AudioFileCreateWithURL`
  - toll free cast between `NSURL` and `CFURLRef`
- write samples in loop
  - number of samples to write per loop
  - number of samples in each wave
- must be careful to handle need for swapping endianness
  - macs are on Intel little-endian CPUs
  - ARM on iPhone is little-endian
  - use `CFSwapInt16HostToBig()`
- use `AudioFileWriteBytes()` to write (for CBR formats only)
  - `AudioFileID`
  - caching flag
  - offset in output to write to
  - number of bytes being written
- must use `AudioFileWritePackets()` for VBR formats
- shows how to rewrite for square, triangle and sawtooth waves
- iOS 4 only handles integers (not floats) for PCM

## Buffers

---

- the time it takes for the audio hardware to produce or consume a packet of audio data is much less than the time it takes to move the packets in and out of memory
- this slowdown is called the von Neumann bottleneck
- use buffers to move data between sources and sinks

- buffering adds complexity and delays audio hardware responsiveness
  - latency
- buffers and latency require a balancing act
  - big buffers cause higher latency
  - lower latency can cause premature buffer drain which results in silence (dropouts) or noise

## Audio Formats

---

- audio data formats and file formats are distinct considerations
  - mp3
  - AIFF handles big-endian PCM
  - WAV handles little-endian PCM
  - MP4 can handle AAC, PCM, and AC3
  - .caf (Core Audio Format) is most content agnostic format
    - any audio format supported by Core Audio
    - mp3
    - AAC
    - Apple Lossless
    - etc
    - good for internal format for an application
    - improves performance

## Ch. 3 Audio Processing with Core Audio

---

### Audio Data Formats

---

- Core Audio views audio data as streams of packets
- `AudioStreamBasicDescription` is critical to most Core Audio programs

```
struct AudioStreamBasicDescription {
    Float64 mSampleRate;
    UInt32 mFormatID;
    UInt32 mFormatFlags;
    UInt32 mBytesPerPacket;
    UInt32 mFramesPerPacket;
    UInt32 mBytesPerFrame;
    UInt32 mChannelsPerFrame;
    UInt32 mBitsPerChannel;
    UInt32 mReserved;
```

```
};  
typedef struct AudioStreamBasicDescription AudioStreamBasicDescription;
```

- `mSampleRate` is defined to be the number of samples per channel per second of uncompressed data
  - makes it equal to the number of frames per second, which means can multiply it by `mFramesPerPacket` to determine packet length in seconds
- `mFormatID` is a four-character code that serves as the name of the format
  - without naming a format, the struct is meaningless, so this value must be defined
  - default types Core Audio supports are defined in `CoreAudioTypes.h`
  - contains the four character codes for the default formats and constants that can be used
  - for example, the `kAudioFormatLinearPCM` format constant used previously chapter is the four-character code `'lpcm'`
- `mFormatFlags` is the format's subtype, or preference panel
  - `UInt32` bit field in which set or check various 1-bit flags
  - setting flags answers questions left open by formats that support non-interleaved data, multiple sample types, or variable data structures
  - the only way to know how to interpret this value is to look it up the documentation — search for “AudioStreamBasicDescription Flags” to find the defined flags
  - formats that do not have flags can set this value to 0
- `mBytesPerPacket` sets details of the structure of the format's data
  - amount/size of data
  - NOTE: VBR formats, by definition, cannot answer this question
  - these formats set this value to 0 and use an `AudioStreamPacketDescription`
- `mFramesPerPacket` subdivides the raw bytes of the packet into some number of frames
  - only for compressed formats - uncompressed formats always set this to 1
  - Core Audio can also support formats with a variable frame rate
  - these formats set this value to 0 and use `AudioStreamPacketDescriptions`
- `mBytesPerFrame` establishes the size of a single frame
  - the total number of digits used to represent each moment in time
  - when a frame does not contain a sample per channel, as with compression, this value is set to 0
- `mChannelsPerFrame` subdivides the frame into channels, regardless of compression
  - this value cannot be 0 because an audio stream with no channels is, by definition, empty

- `mBitsPerChannel` is the sample's bit depth
  - as with bytes per frame, it represents the actual structure of the data
  - compressed formats whose frames do not contain a sample per channel set this value to 0
- `mReserved` member is for data alignment purposes, padding the structure to an even multiple of 8
  - it must always be 0
- the specifics of a particular ASBD are specific to the format being described, i.e. they are an implementation detail
  - do not compare specific fields between formats (unless pertinent)
- in many cases, Core Audio automatically populates the fields of an ASBD
- for compressed data, whose packet structure cannot be adequately derived from its measurements, each packet must be accompanied by an `AudioStreamPacketDescription`

```
struct AudioStreamPacketDescription {
    SInt64 mStartOffset;
    UInt32 mVariableFramesInPacket;
    UInt32 mDataByteSize;
};
typedef struct AudioStreamPacketDescription AudioStreamPacketDescription;
```

- `mStartOffset` represents the packet's location, relative to other packets in the buffer, in bytes
  - important because packets of different sizes cannot use an implied x-axis value as in a linear PCM format
  - cannot find a sample by doing  $\text{Offset} = \text{Time} \times \text{Sample rate} \times \text{Frame size} + \text{Channel number}$
- `mVariableFramesInPacket` represents the number of frames in the packet, but only if the packets use a variable frame rate
  - if the `AudioStreamBasicDescription` has a value for `mFramesPerPacket`, `mVariableFramesInPacket` should be 0, and vice versa
- `mDataByteSize` contains the packet's actual size, in bytes
- many helper API's use additional magic cookies to describe audio data
  - Audio File Services
  - Audio File Stream Services
  - Audio Conversion Services
  - Audio Queue Services

- more
- cookie
  - opaque block of data with contents specific to format being encoded
- when opening file or network stream of compressed data
  - check for magic cookie property
  - if present, read in as block of untyped data and pass to Core Audio without modifying or checking

## Example: Figuring out Formats

---

- to use the property called `kAudioFileGlobalInfo_AvailableStreamDescriptionsForFormat` pass Core Audio a structure called `AudioFileTypeAndFormatID`
  - this structure has two members, a file type and a data format
  - can set both with Core Audio constants found in the documentation or the `AudioFile.h` and `AudioFormat.h` headers
- use an AIFF file with PCM to start
- prepare an `OSStatus` to receive result codes from your Core Audio calls
  - also prepare a `UInt32` to hold the size of the info - have to negotiate before actually retrieving info
- getting a global info property requires a query in advance for the size of the property and to store the size in a pointer to a `UInt32`
- global info calls take a specifier, which acts like an argument to the property call and depends on the property requested
  - in the case of `kAudioFileGlobalInfo_AvailableStreamDescriptionsForFormat` provide the `AudioFileTypeAndFormatID`.
- `AudioFileGetGlobalInfoSize()` calls return amount of data to be received when the global property is retrieved
  - malloc memory to hold the property
- call `AudioFileGetGlobalInfo()` to get the `kAudioFileGlobalInfo_AvailableStreamDescriptionsForFormat`
  - `kAudioFileGlobalInfo_AvailableStreamDescriptionsForFormat`
  - pass in the `AudioFileTypeAndFormatID` and the size of the buffer and a pointer to the buffer
- docs specify that the property call provides an array of `AudioStreamBasicDescriptions`,
  - can determine length of the array by dividing the data size by the size of an ASBD
  - enables a for loop to investigate the ASBDs
- docs stated that the three ASBD fields that get filled in are `mFormatID`, `mFormatFlags`, and `mBitsPerChannel`
  - useful to log the format ID
  - to make it legible, convert it out of the four-character code numeric format and into a readable four-character string



- do this with an endian swap because the UInt32 representation will reorder the bits from their original pseudostring representation
- to pretty print the mFormatId's endian-swapped representation can use the format string %4.4s to force NSLog (or printf) to treat the pointer as an array of 8-bit characters that is exactly four characters long
- mFormatFlags and mBitsPerChannel members are a bit field and numeric value, so print them as ints
- free() the malloc()'d memory to hold the ASBD array when done with it
- book provides additional information about results of code for various formats [here](#)

## Canonical Formats

---

- AAC good for audio fidelity (particularly for music) at high compression rates
- for maximum fidelity can use Apple Lossless which reproduces its source audio perfectly
- the iLBC (Internet Low-Bandwidth Codec) is optimized for speech over potentially unreliable Internet connections
  - makes it well suited for VoIP or in-game chat purposes
- none of above formats are appropriate for editing
- editing better served by the low CPU overhead and losslessness of PCM
  - must have enough RAM and disk space
- any value not specifically set in ASBD defaults to canonical format value
- each platform has two canonical formats
  - AudioSampleType
    - used for I/O situations
  - AudioUnitSampleType
    - introduced in Snow Leopard and iOS
    - used in audio units and for digital signal processing
- on Mac OS X
  - AudioSampleType is 32-bit float
  - AudioUnitSampleType is also a 32-bit float
    - channels must be noninterleaved
- can find the mFormatFlags definitions in the CoreAudioTypes.h header file in the CoreAudio.framework

```
kAudioFormatFlagsCanonical = kAudioFormatFlagIsFloat
                             | kAudioFormatFlagsNativeEndian
                             | kAudioFormatFlagIsPacked
kAudioFormatFlagsAudioUnitCanonical = kAudioFormatFlagIsFloat
                                       | kAudioFormatFlagsNativeEndian
                                       | kAudioFormatFlagIsPacked
                                       | kAudioFormatFlagIsNonInterleaved
```

- on iOS
  - AudioSampleType is 16-bit integer

- `AudioUnitSampleType` is an 8.24-bit fixed-point number
  - 8 bits to the left of the radix point and 24 bits to the right

```
kAudioFormatFlagsCanonical = kAudioFormatFlagIsSignedInteger
                             | kAudioFormatFlagsNativeEndian
                             | kAudioFormatFlagIsPacked
kAudioFormatFlagsAudioUnitCanonical = kAudioFormatFlagIsSignedInteger
                                       | kAudioFormatFlagsNativeEndian
                                       | kAudioFormatFlagIsPacked
                                       | kAudioFormatFlagIsNonInterleaved
                                       | (kAudioUnitSampleFractionBits << kLine
```

- when supplying samples directly to the audio engines (Audio Units and OpenAL) directly, it's advantageous to use the canonical formats
  - saves some data conversions and allows CPU cycles to be used elsewhere

## Processing Audio with Audio Units

---

- core Audio does most of its work at the Audio Units level
  - Audio Queue and OpenAL engines are implemented atop audio units
- each audio unit processes a buffer of samples in some specific way
  - for example, one captures audio from a mic
  - next one downstream may perform an effect on those samples
  - next one may mix it with another source
- analogous to the patch architecture common in the sound industry
- Audio units are connected by audio streams
  - the way audio equipment is patched together with cables
  - audio units have elements, which may have input scope or output scope to indicate that they either accept or produce data
- to connect two units, you set a property connecting an input element of one unit to the output of another
- types of audio units
  - effect units
    - perform digital signal processing (change the audio data in some way)
    - analogous to hardware effects boxes and outboard signal processors
  - instrument units
    - generate audio data representing musical notes, typically from MIDI input, which can itself come from a musical instrument or a software synthesizer
  - generator units
    - also generate audio data, but not from a midi source
    - some units simply generate a signal programmatically, whereas others load data from a network stream or audio file
  - i/o units
    - provide interfaces to input or output hardware, such as a microphone or a speaker

- these are typically implemented on the hardware abstraction layer (HAL)
  - intermediate layer between Core Audio and I/O Kit and the drivers
- converter units
  - reformat audio data back and forth between the canonical formats and other formats
  - these can also merge and split streams, and alter timing and pitch
- mixer units
  - combine audio tracks
  - there are also splitter units that provide multiple outputs from a single input
- panner units
  - use stereo mixing to create panning effects
- offline effect units
  - perform operations on audio data that cannot be done in real time.

## The Pull Model

---

- framework dictates when audio will be provided to be processed rather than programmer
- in a chain of audio units, the last one (probably the I/O unit that will send audio to the speakers or headphones) pulls from the units connected to its input elements
  - each unit call the units upstream from it
- callback = proc in Core Audio
- read proc - callback that receives data
- write proc - callback that produces data
- set a callback property on the audio unit

## Ch. 4 Recording

---

- Audio Queue Services is the highest level playback and recording API in Core Audio
- conveniences
  - unlike with OpenAL and Audio Units, can use encoded formats such as AAC and MP3 with audio queues
  - by default, Audio Queues call back on their own thread, which isolates application programmer from some timing challenges with Audio Units

## All About Audio Queues

---

- simplifies audio playback and recording
  - reduces complexity of audio data management and codecs
  - reduces complexity of underlying hardware interaction
- audio queue is a software interface to piece of hardware
  - transducer (speaker or microphone)

- queue is filled with data and passed to callback which sends audio data somewhere else and places buffer at the back of the queue
- see Figure 4.1 for graphic

## Building a Recorder

---

- audio queue usage is more akin to building a recorder than using a recorder
- provide a callback function that will be called to provide application with buffers of audio captured from the input device (in the recording case) or to fill a buffer for playback
- `AudioQueueNewInput()` takes the following
  - format to record to
  - structure representing a callback function
  - pointer to "user data," which is provided to the callback
  - Core Foundation run loop to use for the callbacks
  - Core Foundation run loop "mode" for callbacks
  - "Flags" that must be set to 0
  - pointer to receive a newly created `AudioQueueRef`

## A `CheckError()` Function

---

- see modified `check_error()` function

## Creating and Using the Audio Queue

---

- create a user info struct for recording Audio Queue callbacks
  - similar to creating a class to encapsulate application logic in C++, etc
- create custom struct and ASBD in main function
- set format of ASBD for Audio Queue
  - `mFormatID - kAudioFormatMPEG4AAC`
  - `mChannelsPerFrame - 2`
- add convenience function to determine correct sample rate for selected input device
- for formats other than PCM, fill in properties of ASBD that can be filled in and allow Core Audio to handle the rest
- check error on output value
- use `AudioQueueNewInput()` to create the audio queue
- after creating the queue, can retrieve a more complete ASBD with `AudioQueueGetProperty` and `kAudioConverterCurrentOutputStreamDescription`
- use info to create audio file to write captured data to with `CFURLCreateWithFileSystemPath()`
- queue returns magic cookie (AAC is a format that uses magic cookies)
  - retrieve from audio queue and set on audio file
  - add convenience function to handle this

- with PCM can calculate buffer sizes to use but cannot with a compressed format
  - $44100 \text{ samples/second} \times 2 \text{ channels} \times 2 \text{ bytes/channel} \times 1 \text{ seconds}$ 
    - requires 176400 bytes to hold a second of 16-bit stereo PCM at 44.1 KHz
- create convenience function to retrieve buffer size from audio queue
- common practice in Core Audio to use 3 buffers
  - one is filled
  - one is drained
  - one is in buffer queue as a spare to account for potential lag
- allocate and enqueue buffers in set up
  - `AudioQueueAllocateBuffer()`
  - `AudioQueueEnqueueBuffer()`
- start the audio queue with `AudioQueueStart()`
- handle user stop
- call `AudioQueueStop()` to stop audio queue
- call magic cookie convenience function for output to file
- clean up audio queue and audio file with `AudioQueueDispose()` and `AudioFileClose()`

## Utility Functions for the Audio Queue

---

- get default device sample rate
  - set up
    - `AudioDeviceID`
    - `AudioObjectPropertyAddress`
      - `mSelector - kAudioHardwarePropertyDefaultInputDevice`
      - `mScope - kAudioObjectPropertyScopeGlobal`
      - `mElement - 0`
  - call `AudioHardwareServiceGetPropertyData()`
  - get the returned input device's sample rate using the property address and calling `AudioHardwareServiceGetPropertyData`
- copying magic cookie from audio queue to audio file
  - get property size of `kAudioConverterCompressionMagicCookie` with `AudioQueueGetPropertySize`
  - get property (`kAudioQueueProperty_MagicCookie`) with `AudioQueueGetProperty()`
  - set property (`kAudioFilePropertyMagicCookieData`) with `AudioFileSetProperty()` on the audio file
  - if size is 0, no cookie exists and can exit, otherwise need to allocate a byte buffer to hold the cookie
- computing recording buffer size for an ASBD
  - first retrieve number of frames (one sample for every channel) in each buffer
    - get this by multiplying the sample rate by the buffer duration

- if the ASBD already has an `mBytesPerFrame` value, as in the case for constant bit rate formats such as PCM, can trivially get the needed byte count by multiplying `mBytesPerFrame` by the frame count.
- if the ASBD already has an `mBytesPerFrame` value, must work at the packet level
  - easy case for this is a constant packet size, indicated by a nonzero `mBytesPerPacket`
  - for hard case, get the audio queue property `kAudioConverterPropertyMaximumOutputPacketSize`, which gives an upper bound
  - in both cases, `maxPacketSize` will have been retrieved as a result
- ASBD might provide a `mFramesPerPacket` value
  - in this case, to determine number of packets, divide the frame count by `mFramesPerPacket` to get a packet count (packets).
  - otherwise, assume the worst case of one frame per packet
- with a frames-per-packet value (which is forced to nonzero for safety) and a maximum size per packet, can multiply the two to get a maximum buffer size

## The Recording Audio Queue Callback

---

- called every time the queue fills one of the buffers with freshly captured audio data
- callback function performs data processing with the data
- start by casting user info `void *` back to proper user info type
- for `AudioFileWritePackets()`, need
  - a file (in the user data struct)
  - a boolean indicating whether to cache the data written
    - false in this case
  - the size of the data buffer to write retrieved from the `inBuffer` parameter's `mAudioDataByteSize`
  - packet descriptions, provided by the callback's `inPacketDesc` parameter
  - an index to which packet in the file to write, which is a running count tracked in recorder's `recordPacket` field
  - number of packets to write, provided by the callback's `inNumPackets` parameter
  - pointer to the audio data, which is the `inBuffer` `mAudioData` pointer
- re-enqueue used buffers

## Ch. 5 Playback

---

- playback requires another form of audio queue

# Defining the Playback Application

---

- create a playback audio queue with `AudioQueueNewOutput()`
  - `AudioStreamBasicDescription`
    - describing the audio format being provided to the queue
  - `AudioQueueOutputCallback`
    - function pointer to a callback function you will write
  - user data pointer to provide to the callback
  - Core Foundation run loop on which to call the callback
  - Core Foundation run loop mode for the callback
  - unused flags parameter that must always be 0
  - pointer to receive the created `AudioQueueRef`
- function pointer to the callback
  - user data pointer
  - queue that is performing the callback
  - `AudioQueueBufferRef` to fill with data

## Setting Up a File-Playing Audio Queue

---

- callback gets void \* to user data struct
  - reads from audio file and passes packets to the queue
- types in user info struct
  - `AudioFileID` playbackFile
  - `SInt64` packetPosition
  - `UInt32` numPacketsToRead
  - Boolean isDone
- define var for the file to read (`CFSTR`)
- use `CFURLCreateWithFileSystemPath()`
- set up and retrieve ASBD using `AudioFileGetProperty()` and `kAudioFilePropertyDataFormat`
- create new Audio Queue for output using `AudioQueueNewOutput()`
- packet - collection of frames which in turn are a collection of samples
- setting up the playback buffers
  - requires working with packets
  - for VBR (mp3 and AAC), use an array of `AudioStreamPacketDescriptions` to provide a map of the contents of the audio buffer
  - call `CalculateBytesForTime()` to calculate the playback buffer size and number of packets to read
  - be sure to allocate proper size for packet descriptions array using `malloc`
  - create convenience function to handle the magic cookie
- see code example for allocating and enqueueing the playback buffers

- call `AudioQueueStart()` to start the queue playback and then `CFRunLoopRunInMode()` with `kCFRunLoopDefaultMode`
  - add delay to ensure queue plays out buffered audio
- call `AudioQueueStop()`, `AudioQueueDispose()` and `AudioFileClose()` to clean-up

## Playback Utility Functions

---

- handling the magic cookie
  - see previous chapter and code example
- calculating buffer size and expected packet count
  - get the maximum packet size for the file's encoding type, which is available from the audio file property `kAudioFilePropertyPacketSizeUpperBound`
  - set up two constants as fail-safe values
    - `maxBufferSize` of 64 KB
    - `minBufferSize` of 16 KB
  - if the ASBD tells how many frames are in a packet
    - calculate how many packets elapse in the given number of seconds and multiply that by the maximum packet size to get a sufficiently large buffer
  - if the ASBD does not tell how many frames are in a packet (no `mFramesPerPacket` value)
    - select an arbitrarily "large enough" value, which is the greater of `maxBufferSize` and `maxPacketSize`
  - in the worst case, at least the buffer will be large enough to hold one packet
  - second if-else applies boundary checks
  - if the calculated buffer size (`outBufferSize`) is larger than both the `maxBufferSize` and the `maxPacketSize`, clamp it to the `maxBufferSize`
    - also check to see that it's not smaller than the `minBufferSize`
  - with an `outBufferSize` calculated
    - divide it by the `maxPacketSize` to figure out how many packets can be safely read from the file on each callback

## The Playback Audio Queue Callback

---

- cast back user info void \*
- use `AudioFileReadPackets()`
  - file to read from
  - cache flag
  - pointer to receive the number of bytes actually read
  - pointer to a buffer to hold packet descriptions
  - index of the first packet you want to read
  - pointer to a maximum number of packets to read (which will be replaced by the number of packets actually read when the function returns)
  - pointer to a buffer to receive the audio data



- use `AudioFileReadPackets()` to read packets
- use `AudioFileEnqueueBuffer()` to enqueue packets
- use `AudioQueueStop()` when the end of file is reached

## Features and Limits of Queue-Based Playback

---

- level-metering
  - `kAudioQueueProperty_EnableLevelMetering`
  - `kAudioQueueProperty_CurrentLevelMeter`
  - `kAudioQueueProperty_CurrentLevelMeterDB`
- properties vs. parameters
  - properties are used for set-up and initialization of an audio object and have values of any type
  - parameters represent values that may be of interest to end user and may change during use
    - always floating point numbers
    - `kAudioQueueParam_Volume` for instance
- benefits
  - quiet handling of decompression from encoded formats such as AAC to PCM stream
  - eliminates threading concerns (callbacks are by default performed by one of queue's internal threads)
- drawbacks
  - each buffer adds latency (.5s for each in this example, 1.5s total)
  - user will not hear read sample until all preceding samples are played out
  - using smaller buffers leads to more callback calls, thus more call overhead, etc
  - inherent latency to queues

## Ch. 6 Conversion

---

### The afconvert Utility

---

- use `--help` to see info detailed here

### Using Audio Converter Services

---

- prefixed with `AudioConverter` (can use for documentation lookup)
- setup user info struct
  - `AudioStreamBasicDescription` `inputFormat`
  - `AudioStreamBasicDescription` `outputFormat`
  - `AudioFileID` `inputFile`
  - `AudioFileID` `outputFile`
  - `UInt64` `inputFilePacketIndex`

- UInt64 inputFilePacketCount
- UInt32 inputFilePacketMaxSize;

## Setting Up Files for Conversion

---

- define file path
- create user info struct
- open file with CFURLCreateWithFilePath()
- call AudioFileOpenURL() and check error
- call CFRelease() on input file URL
- get ASBD from input audio file using AudioFileGetProperty()
- get packet count and maximum packet size properties from the input audio file
- define output ASBD
  - mSampleRate - 44100.0
  - mFormatID - kAudioFormatLinearPCM
  - mFormatFlags - kAudioFormatFlagsBigEndian | kAudioFormatFlagsSignedInteger | kAudioFormatFlagsPacked
  - mBytesPerPacket - 4
  - mFramesPerPacket - 1
  - mBytesPerFrame - 4
  - mChannelsPerFrame - 2
  - mBitsPerChannel - 16
- create output file with CFURLCreateWithFilePath()
- call the custom conversion function
- close the files with AudioFileClose()
- conversion function
  - create AudioConverterRef using AudioConverterNew()
  - determining the size of a packet buffers array and packets-per-buffer count for variable bit rate data
    - see code
  - determine packets per buffer for CBR data using output buffer size / size per packet
  - allocate memory for the audio conversion buffer
  - loop to convert and write data
  - prepare an AudioBufferList to receive converted data
  - call AudioConverterFillComplexBuffer()

- write packets of converted data to output file using `AudioFileWritePackets()`
- call `AudioConverterDispose()` to clean-up
- callback function
  - cast user data struct
  - zero audio buffers
  - determine number of packets that can be read from input
  - allocate a buffer to fill and convert
  - read packets into the conversion buffer
  - update the source file position and `AudioBuffer` members with the results of the read

## Converting with Extended Audio File Services

---

- extended services combine Audio File Services and Audio Converter Services into a simplified API
- example follows same general outline as previous example
- user info struct
  - `AudioStreamBasicDescription` outputFormat
  - `ExtAudioFileRef` inputFile
  - `AudioFileID` outputFile
- call `ExtAudioFileOpenURL()` to open input file
- create ASBD for output format and create audio file
- set the client data format property on the input file using `ExtAudioFileSetProperty()`
- call the conversion function and close the file with `ExtAudioFileDispose()` and `AudioFileClose()`
- to read and convert for extended audio files
  - determine the size of the output buffer and packets-per-buffer count
  - allocate a buffer for receiving data from an extended audio file
  - run the read-convert-write loop after setting up an `AudioBuffer`
    - use `ExtAudioFileRead()` to read
    - terminate if no frames are read
    - use `AudioFileWritePackets()` to write
    - advance the write position

## Ch. 7 Audio Units: Generators, Effects, and Rendering

---

### Where the Magic Happens

---

- Audio Units is lowest level of audio functionality that application developers need

- enables functionality that is impractical or impossible with the higher-level APIs
- increases difficulty of use and potential for serious error

## How Audio Units Work

---

- works on streams of audio
- set up units, create connections, tell them to start processing
- analogous to the setup and flow of audio hardware
- types of audio units
  - generator units
    - create a stream of audio from some source, such as files, the network, or memory
  - instrument units
    - similar to generator units, produce a stream of synthesized audio from MIDI data
  - mixer units
    - combine multiple streams into one or more streams. the mix can be performed in two dimensions (for stereo panning) or in a simulated three-dimensional sound field
  - effect units
    - perform some sort of digital signal processing on a stream, usually producing an audible effect such as a reverb, a pitch change, noise filtering, and so on
  - converter units
    - perform transformations that are generally not meant to deliver user-audible effects
    - includes units to convert between different flavors of pcm (to change sample rate or bit depth, for example), adjust playback speed, and so on
  - output units
    - interface with audio input and/ or output hardware, enabling you to capture audio from input devices and play it out to output devices
    - name is misleading because these are potentially input/output units
- works with a pull model - software objects that need data pull samples from other objects using callbacks

## Sizing Up the Audio Units

---

- type, subtype, and manufacturer are in AUComponent.h

### Audio Unit Subtypes for Generator Units (Type `kAudioUnitType_Generator`)

- `kAudioUnitSubType_ScheduledSoundPlayer`
  - schedules audio to be played at a specific time
- `kAudioUnitSubType_AudioFilePlayer`
  - plays audio from a file
- `kAudioUnitSubType_NetReceive`
  - receives network audio from a corresponding `kAudioUnitSubType_NetSend` unit on another host or in another application

## **Audio Unit Subtypes for Instrument Units (Type `kAudioUnitType_MusicDevice`)**

- `kAudioUnitSubType_DLSSynth`
  - multi-timbral music synthesizer that accepts MIDI commands
  - works with DLS or SoundFont formats

## **Audio Unit Subtypes for Mixer Units (Type `kAudioUnitType_Mixer`)**

- `kAudioUnitSubType_MultiChannelMixer`
  - mixes any number of single or multi-channel input buses to one output bus
- `kAudioUnitSubType_StereoMixer`
  - mixes any number of mono or stereo input buses to one stereo output
- `kAudioUnitSubType_3DMixer`
  - mixes any number of mono or stereo input buses
  - mono inputs can be panned with 3D parameters
  - output is one bus of 2 to 8 channels
- `kAudioUnitSubType_MatrixMixer`
  - mixes any number of input and output buses, with any number of channels per bus
  - supports highly configurable mapping of inputs and outputs

## **Audio Unit Subtypes for Panner Units (Type `kAudioUnitType_Panner`)**

- `kAudioUnitSubType_SphericalHeadPanner`
  - uses "spherical head" model to produce stereo output
- `kAudioUnitSubType_VectorPanner`
  - uses pan between adjacent channels in 3D space to create surround output
- `kAudioUnitSubType_SoundFilePanner`
  - uses "sound field" model to produce stereo output
- `kAudioUnitSubType_HRTFPanner`
  - uses head-related transfer function to produce stereo output

## **Audio Unit Subtypes for Effect Units (Type `kAudioUnitType_Effect`)**

- `kAudioUnitSubType_Delay`
  - digital delay effect
- `kAudioUnitSubType_LowPassFilter`
- `kAudioUnitSubType_HighPassFilter`
- `kAudioUnitSubType_BandPassFilter`
- `kAudioUnitSubType_HighShelfFilter`
- `kAudioUnitSubType_LowShelfFilter`
- `kAudioUnitSubType_ParametricEQ`
- `kAudioUnitSubType_GraphicEQ`
- `kAudioUnitSubType_PeakLimiter`
- `kAudioUnitSubType_DynamicsProcessor`

- kAudioUnitSubType\_MultiBandCompressor
- kAudioUnitSubType\_MatrixReverb
- kAudioUnitSubType\_SampleDelay
- kAudioUnitSubType\_Pitch
- kAudioUnitSubType\_AUFilter
  - adjusts gain for five bands of frequencies
- kAudioUnitSubType\_NetSend
  - sends audio across network or between applications
- kAudioUnitSubType\_Distortion
- kAudioUnitSubType\_RogerBeep
  - produces a beep, similar to sound of walkie-talkie release, when the input level drops below a certain threshold for a certain amount of time

### **Audio Unit Subtypes for Converter Units (Type kAudioUnitType\_FormatConverter)**

- kAudioUnitSubType\_AUConverter
  - uses audio converter services to perform LPCM conversions
- kAudioUnitSubType\_Varispeed
  - changes playback speed, pitch-shifts audio as a result
  - faster = higher pitch
- kAudioUnitSubType\_TimePitch
  - similar to varispeed
  - allows for independent control of playback speed and pitch
  - can play faster without changing pitch
- kAudioUnitSubType\_DeferredRenderer
  - pulls input from a thread other than the caller
- kAudioUnitSubType\_Splitter
  - splits one input bus into two identical output buses
- kAudioUnitSubType\_Merger
  - merges two input buses into one output bus

### **Audio Unit Subtypes for Output Units (Type kAudioUnitType\_Output)**

- kAudioUnitSubType\_GenericOutput
  - output unit not tied to audio hardware
  - can be used to perform software rendering of audio in connected units
- kAudioUnitSubType\_SystemOutput
  - output to the device used for alerts and other UI sounds
- kAudioUnitSubType\_DefaultOutput
  - output to the device selected in System Preferences -> Sound
- kAudioUnitSubType\_HALOutput
  - input from and output to any supported audio device