

Ch. 1 Overview of Core Audio

- engine behind sound on mac or iOS
 - C API
 - exposed in Objective-C, C++, and Swift
- high-level when compared to general audio programming environments
- low-level from an application developer standpoint
- very different from Cocoa and other C-based APIs (Quartz, Core Foundation)
- when not to use
 - simple file playback
 - use AppKit and NSSound or AVAudioPlayer from AV Foundation
 - can also use QTKit on mac

The Core Audio Frameworks

- two groups
 - audio engines - audio stream processing
 - helper APIs - facilitate audio IO to and from these frameworks
- mac and iOS audio engine APIs
 - Audio Units
 - Core Audio passes audio buffers to graphs of Audio Units which do work on the buffers
 - Audio Queues
 - abstraction on top of audio units that simplifies playback and recording
 - uses callbacks
 - OpenAL
 - industry standard 3D audio API (designed to resemble OpenGL)
- helper APIs
 - Audio File Services
 - abstracts access to various file types
 - Audio File Stream Services
 - network and other stream sources of audio
 - Audio Converter Services
 - format conversion
 - Extended Audio File Services
 - allows simultaneous reading and writing and conversion
 - Core MIDI
 - MIDI access

- Audio Session Services (iOS-only)
 - coordinates use of resources with system on iOS

Core Audio Conventions

- C function calls
 - brings a number of difficulties that modern programmers may not be equipped to deal with
- does mimic object orientation within C in many ways
- Core Foundation is the main C framework and underlies Foundation
 - many of the Objective-C components of Core Foundation gain their functionality by calling C
- Core Audio is very similar in design to Core Foundation

Your First Core Audio Application

- uses Objective-C in Xcode
- see code for details
- notes
 - need to convert from C strings to NSString/CFStringRef
 - can specify UTF-8
 - can use stringByExpandingTildeInPath
 - must convert file paths to NSURL to use in audio file APIs
 - use local variable reference of type AudioFileID type to refer to audio file objects
 - use OSStatus to signal failure or success through return types
 - should check on all Core Audio calls
 - noErr (0) signals success
 - calls AudioFileOpenURL to open file
 - in Objective-C, NSURL can be toll-free cast to CFURLRef
 - to get metadata, ask for kAudioFilePropertyInfoDictionary
 - requires advance allocation of memory for returned metadata
 - to get size for allocation, call AudioFileGetPropertyInfo with AudioFileID
 - refer to docs and headers for info on the various results of getting properties
 - call AudioFileGetProperty
 - Core Foundation does not offer autorelease so retain/release counts must be managed manually
 - AudioFileID has its own clean-up call, AudioFileClose
 - many Core Audio objects have custom clean-up calls, be aware of these
 - check return codes on Core Audio clean-up calls

Running the Example

- requires linking in frameworks in Xcode
- requires passing command-line arguments
- NOTE: the metadata results will vary wildly depending on files/file types/systems
- NOTE: Core Audio requires an adherence to a large number of implicit contracts between various types of calls
 - BE PATIENT AND WORK WITH THE CODE

Core Audio Properties

- preparing and calling property getters and setters in Core Audio is an essential process
- properties
 - key-value pairs
 - keys are enumerated integers
 - values are types defined by APIs
 - values that are retrieved or set are dependent on the specific property
- for `AudioFileGetProperty`, the keys are 32-bit integers that can be represented as four character codes
 - NOTE: use single-quotes to represent char literals in C
- many Core Audio types have related `*GetProperty()` functions
- function signature adds information based on naming conventions of function parameters
 - in - only used for input to the function
 - out - only used for output from the function
 - io - used for input and output from a function

Ch. 2 The Story of Sound

Making Waves

Digital Audio

DIY Samples

- need to set
 - sample rate
 - length of audio sample to create

- filename/path of output file
- file format
- frequency of audio signal (as command-line arg)
- generate full filename from file format
- need to populate info in `AudioStreamBasicDescription` struct
 - channels
 - format
 - bitrate
 - etc
- initialize remaining fields of `AudioStreamBasicDescription` to 0 (using `memset` in C)
 - Core Audio populates certain fields depending on format
- set info in `AudioStreamBasicDescription`
 - single channel
 - PCM
 - 44100
 - 16-bit
 - 2 bytes per frame (one channel * 2 bytes of sample data)
 - endianness (for PCM)
 - numeric format of samples (`kAudioFormatFlagsSignedInteger`)
 - packed (sample uses all bits available in each byte) vs unpacked
- packets are only useful for VBR formats
 - `bytesPerFrame` and `bytesPerPacket` are equal for CBR formats
- AIFF only handles big endian PCM
- numeric format and packed vs unpacked are flags that must be bitwise or'ed together
- create `AudioFileID` for audio file using `AudioFileCreateWithURL`
 - toll free cast between `NSURL` and `CFURLRef`
- write samples in loop
 - number of samples to write per loop
 - number of samples in each wave
- must be careful to handle need for swapping endianness
 - macs are on Intel little-endian CPUs
 - ARM on iPhone is little-endian
 - use `CFSwapInt16HostToBig()`
- use `AudioFileWriteBytes()` to write (for CBR formats only)
 - `AudioFileID`
 - caching flag
 - offset in output to write to
 - number of bytes being written
- must use `AudioFileWritePackets()` for VBR formats

- shows how to rewrite for square, triangle and sawtooth waves
- iOS 4 only handles integers (not floats) for PCM

Buffers

- the time it takes for the audio hardware to produce or consume a packet of audio data is much less than the time it takes to move the packets in and out of memory
- this slowdown is called the von Neumann bottleneck
- use buffers to move data between sources and sinks
- buffering adds complexity and delays audio hardware responsiveness
 - latency
- buffers and latency require a balancing act
 - big buffers cause higher latency
 - lower latency can cause premature buffer drain which results in silence (dropouts) or noise

Audio Formats

- audio data formats and file formats are distinct considerations
 - mp3
 - AIFF handles big-endian PCM
 - WAV handles little-endian PCM
 - MP4 can handle AAC, PCM, and AC3
 - .caf (Core Audio Format) is most content agnostic format
 - any audio format supported by Core Audio
 - mp3
 - AAC
 - Apple Lossless
 - etc
 - good for internal format for an application
 - improves performance

Ch. 3 Audio Processing with Core Audio

Audio Data Formats

- Core Audio views audio data as streams of packets
- `AudioStreamBasicDescription` is critical to most Core Audio programs

```

struct AudioStreamBasicDescription {
    Float64 mSampleRate;
    UInt32 mFormatID;
    UInt32 mFormatFlags;
    UInt32 mBytesPerPacket;
    UInt32 mFramesPerPacket;
    UInt32 mBytesPerFrame;
    UInt32 mChannelsPerFrame;
    UInt32 mBitsPerChannel;
    UInt32 mReserved;
};
typedef struct AudioStreamBasicDescription AudioStreamBasicDescription;

```

- mSampleRate is defined to be the number of samples per channel per second of uncompressed data
 - makes it equal to the number of frames per second, which means can multiply it by mFramesPerPacket to determine packet length in seconds
- mFormatID is a four-character code that serves as the name of the format
 - without naming a format, the struct is meaningless, so this value must be defined
 - default types Core Audio supports are defined in CoreAudioTypes.h
 - contains the four character codes for the default formats and constants that can be used
 - for example, the kAudioFormatLinearPCM format constant used previously chapter is the four-character code 'lpcm'
- mFormatFlags is the format's subtype, or preference panel
 - UInt32 bit field in which set or check various 1-bit flags
 - setting flags answers questions left open by formats that support non-interleaved data, multiple sample types, or variable data structures
 - the only way to know how to interpret this value is to look it up the documentation — search for “AudioStreamBasicDescription Flags” to find the defined flags
 - formats that do not have flags can set this value to 0
- mBytesPerPacket sets details of the structure of the format's data
 - amount/size of data
 - NOTE: VBR formats, by definition, cannot answer this question
 - these formats set this value to 0 and use an AudioStreamPacketDescription
- mFramesPerPacket subdivides the raw bytes of the packet into some number of frames
 - only for compressed formats - uncompressed formats always set this to 1
 - Core Audio can also support formats with a variable frame rate
 - these formats set this value to 0 and use AudioStreamPacketDescriptions

- `mBytesPerFrame` establishes the size of a single frame
 - the total number of digits used to represent each moment in time
 - when a frame does not contain a sample per channel, as with compression, this value is set to 0
- `mChannelsPerFrame` subdivides the frame into channels, regardless of compression
 - this value cannot be 0 because an audio stream with no channels is, by definition, empty
- `mBitsPerChannel` is the sample's bit depth
 - as with bytes per frame, it represents the actual structure of the data
 - compressed formats whose frames do not contain a sample per channel set this value to 0
- `mReserved` member is for data alignment purposes, padding the structure to an even multiple of 8
 - it must always be 0
- the specifics of a particular ASBD are specific to the format being described, i.e. they are an implementation detail
 - do not compare specific fields between formats (unless pertinent)
- in many cases, Core Audio automatically populates the fields of an ASBD
- for compressed data, whose packet structure cannot be adequately derived from its measurements, each packet must be accompanied by an `AudioStreamPacketDescription`

```
struct AudioStreamPacketDescription {
    SInt64 mStartOffset;
    UInt32 mVariableFramesInPacket;
    UInt32 mDataByteSize;
};
typedef struct AudioStreamPacketDescription AudioStreamPacketDescription;
```

- `mStartOffset` represents the packet's location, relative to other packets in the buffer, in bytes
 - important because packets of different sizes cannot use an implied x-axis value as in a linear PCM format
 - cannot find a sample by doing $\text{Offset} = \text{Time} \times \text{Sample rate} \times \text{Frame size} + \text{Channel number}$

- `mVariableFramesInPacket` represents the number of frames in the packet, but only if the packets use a variable frame rate
 - if the `AudioStreamBasicDescription` has a value for `mFramesPerPacket`, `mVariableFramesInPacket` should be 0, and vice versa
- `mDataByteSize` contains the packet's actual size, in bytes
- many helper API's use additional magic cookies to describe audio data
 - Audio File Services
 - Audio File Stream Services
 - Audio Conversion Services
 - Audio Queue Services
 - more
- cookie
 - opaque block of data with contents specific to format being encoded
- when opening file or network stream of compressed data
 - check for magic cookie property
 - if present, read in as block of untyped data and pass to Core Audio without modifying or checking

Example: Figuring out Formats

- to use the property called `kAudioFileGlobalInfo_AvailableStreamDescriptionsForFormat` pass Core Audio a structure called `AudioFileTypeAndFormatID`
 - this structure has two members, a file type and a data format
 - can set both with Core Audio constants found in the documentation or the `AudioFile.h` and `AudioFormat.h` headers
- use an AIFF file with PCM to start
- prepare an `OSStatus` to receive result codes from your Core Audio calls
 - also prepare a `UInt32` to hold the size of the info - have to negotiate before actually retrieving info
- getting a global info property requires a query in advance for the size of the property and to store the size in a pointer to a `UInt32`
- global info calls take a specifier, which acts like an argument to the property call and depends on the property requested
 - in the case of `kAudioFileGlobalInfo_AvailableStreamDescriptionsForFormat` provide the `AudioFileTypeAndFormatID`.

- `AudioFileGetGlobalInfoSize()` calls return amount of data to be received when the global property is retrieved
 - malloc memory to hold the property
- call `AudioFileGetGlobalInfo()` to get the `kAudioFileGlobalInfo_AvailableStreamDescriptionsForFormat`
 - `kAudioFileGlobalInfo_AvailableStreamDescriptionsForFormat`
 - pass in the `AudioFileTypeAndFormatID` and the size of the buffer and a pointer to the buffer
- docs specify that the property call provides an array of `AudioStreamBasicDescriptions`,
 - can determine length of the array by dividing the data size by the size of an ASBD
 - enables a for loop to investigate the ASBDs
- docs stated that the three ASBD fields that get filled in are `mFormatID`, `mFormatFlags`, and `mBitsPerChannel`
 - useful to log the format ID
 - to make it legible, convert it out of the four-character code numeric format and into a readable four-character string
 - do this with an endian swap because the `UInt32` representation will reorder the bits from their original pseudostring representation
 - to pretty print the `mFormatID`'s endian-swapped representation can use the format string `%4.4s` to force `NSLog` (or `printf`) to treat the pointer as an array of 8-bit characters that is exactly four characters long
 - `mFormatFlags` and `mBitsPerChannel` members are a bit field and numeric value, so print them as ints
- `free()` the `malloc()`'d memory to hold the ASBD array when done with it
- book provides additional information about results of code for various formats [here](#)

Canonical Formats

- AAC good for audio fidelity (particularly for music) at high compression rates
- for maximum fidelity can use Apple Lossless which reproduces its source audio perfectly
- the iLBC (Internet Low-Bandwidth Codec) is optimized for speech over potentially unreliable Internet connections
 - makes it well suited for VoIP or in-game chat purposes
- none of above formats are appropriate for editing
- editing better served by the low CPU overhead and losslessness of PCM
 - must have enough RAM and disk space
- any value not specifically set in ASBD defaults to canonical format value
- each platform has two canonical formats
 - `AudioSampleType`
 - used for I/O situations

- `AudioUnitSampleType`
 - introduced in Snow Leopard and iOS
 - used in audio units and for digital signal processing
- on Mac OS X
 - `AudioSampleType` is 32-bit float
 - `AudioUnitSampleType` is also a 32-bit float
 - channels must be noninterleaved
- can find the `mFormatFlags` definitions in the `CoreAudioTypes.h` header file in the `CoreAudio.framework`

```
kAudioFormatFlagsCanonical = kAudioFormatFlagIsFloat
                            | kAudioFormatFlagsNativeEndian
                            | kAudioFormatFlagIsPacked
kAudioFormatFlagsAudioUnitCanonical = kAudioFormatFlagIsFloat
                                       | kAudioFormatFlagsNativeEndian
                                       | kAudioFormatFlagIsPacked
                                       | kAudioFormatFlagIsNonInterleaved
```

- on iOS
 - `AudioSampleType` is 16-bit integer
 - `AudioUnitSampleType` is an 8.24-bit fixed-point number
 - 8 bits to the left of the radix point and 24 bits to the right

```
kAudioFormatFlagsCanonical = kAudioFormatFlagIsSignedInteger
                            | kAudioFormatFlagsNativeEndian
                            | kAudioFormatFlagIsPacked
kAudioFormatFlagsAudioUnitCanonical = kAudioFormatFlagIsSignedInteger
                                       | kAudioFormatFlagsNativeEndian
                                       | kAudioFormatFlagIsPacked
                                       | kAudioFormatFlagIsNonInterleaved
                                       | (kAudioUnitSampleFractionBits << kLine
```

- when supplying samples directly to the audio engines (Audio Units and OpenAL) directly, it's advantageous to use the canonical formats
 - saves some data conversions and allows CPU cycles to be used elsewhere

Processing Audio with Audio Units

- core Audio does most of its work at the Audio Units level
 - Audio Queue and OpenAL engines are implemented atop audio units
- each audio unit processes a buffer of samples in some specific way
 - for example, one captures audio from a mic
 - next one downstream may perform an effect on those samples
 - next one may mix it with another source

- analogous to the patch architecture common in the sound industry
- Audio units are connected by audio streams
 - the way audio equipment is patched together with cables
 - audio units have elements, which may have input scope or output scope to indicate that they either accept or produce data
- to connect two units, you set a property connecting an input element of one unit to the output of another
- types of audio units
 - effect units
 - perform digital signal processing (change the audio data in some way)
 - analogous to hardware effects boxes and outboard signal processors
 - instrument units
 - generate audio data representing musical notes, typically from MIDI input, which can itself come from a musical instrument or a software synthesizer
 - generator units
 - also generate audio data, but not from a midi source
 - some units simply generate a signal programmatically, whereas others load data from a network stream or audio file
 - i/ o units
 - provide interfaces to input or output hardware, such as a microphone or a speaker
 - these are typically implemented on the hardware abstraction layer (HAL)
 - intermediate layer between Core Audio and I/O Kit and the drivers
 - converter units
 - reformat audio data back and forth between the canonical formats and other formats
 - these can also merge and split streams, and alter timing and pitch
 - mixer units
 - combine audio tracks
 - there are also splitter units that provide multiple outputs from a single input
 - panner units
 - use stereo mixing to create panning effects
 - offline effect units
 - perform operations on audio data that cannot be done in real time.

The Pull Model

- framework dictates when audio will be provided to be processed rather than programmer
- in a chain of audio units, the last one (probably the I/O unit that will send audio to the speakers or headphones) pulls from the units connected to its input elements
 - each unit call the units upstream from it
- callback = proc in Core Audio

- read proc - callback that receives data
- write proc - callback that produces data
- set a callback property on the audio unit

Ch. 4 Recording

- Audio Queue Services is the highest level playback and recording API in Core Audio
- conveniences
 - unlike with OpenAL and Audio Units, can use encoded formats such as AAC and MP3 with audio queues
 - by default, Audio Queues call back on their own thread, which isolates application programmer from some timing challenges with Audio Units

All About Audio Queues

- simplifies audio playback and recording
 - reduces complexity of audio data management and codecs
 - reduces complexity of underlying hardware interaction
- audio queue is a software interface to piece of hardware
 - transducer (speaker or microphone)
- queue is filled with data and passed to callback which sends audio data somewhere else and places buffer at the back of the queue
- see Figure 4.1 for graphic

Building a Recorder

- audio queue usage is more akin to building a recorder than using a recorder
- provide a callback function that will be called to provide application with buffers of audio captured from the input device (in the recording case) or to fill a buffer for playback
- `AudioQueueNewInput()` takes the following
 - format to record to
 - structure representing a callback function
 - pointer to "user data," which is provided to the callback
 - Core Foundation run loop to use for the callbacks
 - Core Foundation run loop "mode" for callbacks
 - "Flags" that must be set to 0
 - pointer to receive a newly created `AudioQueueRef`

A CheckError() Function

- see modified check_error() function

Creating and Using the Audio Queue

- create a user info struct for recording Audio Queue callbacks
 - similar to creating a class to encapsulate application logic in C++, etc
- create custom struct and ASBD in main function
- set format of ASBD for Audio Queue
 - mFormatID - kAudioFormatMPEG4AAC
 - mChannelsPerFrame - 2
- add convenience function to determine correct sample rate for selected input device
- for formats other than PCM, fill in properties of ASBD that can be filled in and allow Core Audio to handle the rest
- check error on output value
- use AudioQueueNewInput() to create the audio queue
- after creating the queue, can retrieve a more complete ASBD with AudioQueueGetProperty and kAudioConverterCurrentOutputStreamDescription
- use info to create audio file to write captured data to with CFURLCreateWithFileSystemPath()
- queue returns magic cookie (AAC is a format that uses magic cookies)
 - retrieve from audio queue and set on audio file
 - add convenience function to handle this
- with PCM can calculate buffer sizes to use but cannot with a compressed format
 - $44100 \text{ samples/second} \times 2 \text{ channels} \times 2 \text{ bytes/channel} \times 1 \text{ seconds}$
 - requires 176400 bytes to hold a second of 16-bit stereo PCM at 44.1 KHz
- create convenience function to retrieve buffer size from audio queue
- common practice in Core Audio to use 3 buffers
 - one is filled
 - one is drained
 - one is in buffer queue as a spare to account for potential lag
- allocate and enqueue buffers in set up
 - AudioQueueAllocateBuffer()
 - AudioQueueEnqueueBuffer()
- start the audio queue with AudioQueueStart()
- handle user stop
- call AudioQueueStop() to stop audio queue
- call magic cookie convenience function for output to file

- clean up audio queue and audio file with `AudioQueueDispose()` and `AudioFileClose()`

Utility Functions for the Audio Queue

- get default device sample rate
 - set up
 - `AudioDeviceID`
 - `AudioObjectPropertyAddress`
 - `mSelector - kAudioHardwarePropertyDefaultInputDevice`
 - `mScope - kAudioObjectPropertyScopeGlobal`
 - `mElement - 0`
 - call `AudioHardwareServiceGetPropertyData()`
 - get the returned input device's sample rate using the property address and calling `AudioHardwareServiceGetPropertyData`
- copying magic cookie from audio queue to audio file
 - get property size of `kAudioConverterCompressionMagicCookie` with `AudioQueueGetPropertySize`
 - get property (`kAudioQueueProperty_MagicCookie`) with `AudioQueueGetProperty()`
 - set property (`kAudioFilePropertyMagicCookieData`) with `AudioFileSetProperty()` on the audio file
 - if size is 0, no cookie exists and can exit, otherwise need to allocate a byte buffer to hold the cookie
- computing recording buffer size for an ASBD
 - first retrieve number of frames (one sample for every channel) in each buffer
 - get this by multiplying the sample rate by the buffer duration
 - if the ASBD already has an `mBytesPerFrame` value, as in the case for constant bit rate formats such as PCM, can trivially get the needed byte count by multiplying `mBytesPerFrame` by the frame count.
 - if the ASBD already has an `mBytesPerFrame` value, must work at the packet level
 - easy case for this is a constant packet size, indicated by a nonzero `mBytesPerPacket`
 - for hard case, get the audio queue property `kAudioConverterPropertyMaximumOutputPacketSize`, which gives an upper bound
 - in both cases, `maxPacketSize` will have been retrieved as a result
 - ASBD might provide a `mFramesPerPacket` value
 - in this case, to determine number of packets, divide the frame count by `mFramesPerPacket` to get a packet count (packets).
 - otherwise, assume the worst case of one frame per packet
 - with a frames-per-packet value (which is forced to nonzero for safety) and a maximum size per packet, can multiply the two to get a maximum buffer size

The Recording Audio Queue Callback

- called every time the queue fills one of the buffers with freshly captured audio data
- callback function performs data processing with the data
- start by casting user info void * back to proper user info type
- for AudioFileWritePackets(), need
 - a file (in the user data struct)
 - a boolean indicating whether to cache the data written
 - false in this case
 - the size of the data buffer to write retrieved from the inBuffer parameter's mAudioDataByteSize
 - packet descriptions, provided by the callback's inPacketDesc parameter
 - an index to which packet in the file to write, which is a running count tracked in recorder's recordPacket field
 - number of packets to write, provided by the callback's inNumPackets parameter
 - pointer to the audio data, which is the inBuffer mAudioData pointer
- re-enqueue used buffers

Ch. 5 Playback

- playback requires another form of audio queue

Defining the Playback Application

- create a playback audio queue with AudioQueueNewOutput()
 - AudioStreamBasicDescription
 - describing the audio format being provided to the queue
 - AudioQueueOutputCallback
 - function pointer to a callback function you will write
 - user data pointer to provide to the callback
 - Core Foundation run loop on which to call the callback
 - Core Foundation run loop mode for the callback
 - unused flags parameter that must always be 0
 - pointer to receive the created AudioQueueRef
- function pointer to the callback
 - user data pointer
 - queue that is performing the callback
 - AudioQueueBufferRef to fill with data

Setting Up a File-Playing Audio Queue

- callback gets void * to user data struct
 - reads from audio file and passes packets to the queue
- types in user info struct
 - AudioFileID playbackFile
 - SInt64 packetPosition
 - UInt32 numPacketsToRead
 - Boolean isDone
- define var for the file to read (CFSTR)
- use CFURLCreateWithFileSystemPath()
- set up and retrieve ASBD using AudioFileGetProperty() and kAudioFilePropertyDataFormat
- create new Audio Queue for output using AudioQueueNewOutput()
- packet - collection of frames which in turn are a collection of samples
- setting up the playback buffers
 - requires working with packets
 - for VBR (mp3 and AAC), use an array of AudioStreamPacketDescriptions to provide a map of the contents of the audio buffer
 - call CalculateBytesForTime() to calculate the playback buffer size and number of packets to read
 - be sure to allocate proper size for packet descriptions array using malloc
 - create convenience function to handle the magic cookie
- see code example for allocating and enqueueing the playback buffers
- call AudioQueueStart() to start the queue playback and then CFRunLoopRunInMode() with kCFRunLoopDefaultMode
 - add delay to ensure queue plays out buffered audio
- call AudioQueueStop(), AudioQueueDispose() and AudioFileClose() to clean-up

Playback Utility Functions

- handling the magic cookie
 - see previous chapter and code example
- calculating buffer size and expected packet count
 - get the maximum packet size for the file's encoding type, which is available from the audio file property kAudioFilePropertyPacketSizeUpperBound
 - set up two constants as fail-safe values
 - maxBufferSize of 64 KB
 - minBufferSize of 16 KB

- if the ASBD tells how many frames are in a packet
 - calculate how many packets elapse in the given number of seconds and multiply that by the maximum packet size to get a sufficiently large buffer
- if the ASBD does not tell how many frames are in a packet (no `mFramesPerPacket` value)
 - select an arbitrarily “large enough” value, which is the greater of `maxBufferSize` and `maxPacketSize`
- in the worst case, at least the buffer will be large enough to hold one packet
- second if-else applies boundary checks
- if the calculated buffer size (`outBufferSize`) is larger than both the `maxBufferSize` and the `maxPacketSize`, clamp it to the `maxBufferSize`
 - also check to see that it’s not smaller than the `minBufferSize`
- with an `outBufferSize` calculated
 - divide it by the `maxPacketSize` to figure out how many packets can be safely read from the file on each callback

The Playback Audio Queue Callback

- cast back user info void *
- use `AudioFileReadPackets()`
 - file to read from
 - cache flag
 - pointer to receive the number of bytes actually read
 - pointer to a buffer to hold packet descriptions
 - index of the first packet you want to read
 - pointer to a maximum number of packets to read (which will be replaced by the number of packets actually read when the function returns)
 - pointer to a buffer to receive the audio data
- use `AudioFileReadPackets()` to read packets
- use `AudioFileEnqueueBuffer()` to enqueue packets
- use `AudioQueueStop()` when the end of file is reached

Features and Limits of Queue-Based Playback

- level-metering
 - `kAudioQueueProperty_EnableLevelMetering`
 - `kAudioQueueProperty_CurrentLevelMeter`
 - `kAudioQueueProperty_CurrentLevelMeterDB`

- properties vs. parameters
 - properties are used for set-up and initialization of an audio object and have values of any type
 - parameters represent values that may be of interest to end user and may change during use
 - always floating point numbers
 - kAudioQueueParam_Volume for instance
- benefits
 - quiet handling of decompression from encoded formats such as AAC to PCM stream
 - eliminates threading concerns (callbacks are by default performed by one of queue's internal threads)
- drawbacks
 - each buffer adds latency (.5s for each in this example, 1.5s total)
 - user will not hear read sample until all preceding samples are played out
 - using smaller buffers leads to more callback calls, thus more call overhead, etc
 - inherent latency to queues

Ch. 6 Conversion

The afconvert Utility

- use `--help` to see info detailed here

Using Audio Converter Services

- prefixed with AudioConverter (can use for documentation lookup)
- setup user info struct
 - AudioStreamBasicDescription inputFormat
 - AudioStreamBasicDescription outputFormat
 - AudioFileID inputFile
 - AudioFileID outputFile
 - UInt64 inputFilePacketIndex
 - UInt64 inputFilePacketCount
 - UInt32 inputFilePacketMaxSize;

Setting Up Files for Conversion

- define file path
- create user info struct

- open file with `CFURLCreateWithFilePath()`
- call `AudioFileOpenURL()` and check error
- call `CFRelease()` on input file URL
- get ASBD from input audio file using `AudioFileGetProperty()`
- get packet count and maximum packet size properties from the input audio file
- define output ASBD
 - `mSampleRate` - 44100.0
 - `mFormatID` - `kAudioFormatLinearPCM`
 - `mFormatFlags` - `kAudioFormatFlagsBigEndian` | `kAudioFormatFlagsSignedInteger` | `kAudioFormatFlagsPacked`
 - `mBytesPerPacket` - 4
 - `mFramesPerPacket` - 1
 - `mBytesPerFrame` - 4
 - `mChannelsPerFrame` - 2
 - `mBitsPerChannel` - 16
- create output file with `CFURLCreateWithFilePath()`
- call the custom conversion function
- close the files with `AudioFileClose()`
- conversion function
 - create `AudioConverterRef` using `AudioConverterNew()`
 - determining the size of a packet buffers array and packets-per-buffer count for variable bit rate data
 - see code
 - determine packets per buffer for CBR data using output buffer size / size per packet
 - allocate memory for the audio conversion buffer
 - loop to convert and write data
 - prepare an `AudioBufferList` to receive converted data
 - call `AudioConverterFillComplexBuffer()`
 - write packets of converted data to output file using `AudioFileWritePackets()`
 - call `AudioConverterDispose()` to clean-up
- callback function
 - cast user data struct
 - zero audio buffers

- determine number of packets that can be read from input
- allocate a buffer to fill and convert
- read packets into the conversion buffer
- update the source file position and AudioBuffer members with the results of the read

Converting with Extended Audio File Services

- extended services combine Audio File Services and Audio Converter Services into a simplified API
- example follows same general outline as previous example
- user info struct
 - AudioStreamBasicDescription outputFormat
 - ExtAudioFileRef inputFile
 - AudioFileID outputFile
- call ExtAudioFileOpenURL() to open input file
- create ASBD for output format and create audio file
- set the client data format property on the input file using ExtAudioFileSetProperty()
- call the conversion function and close the file with ExtAudioFileDispose() and AudioFileClose()
- to read and convert for extended audio files
 - determine the size of the output buffer and packets-per-buffer count
 - allocate a buffer for receiving data from an extended audio file
 - run the read-convert-write loop after setting up an AudioBuffer
 - use ExtAudioFileRead() to read
 - terminate if no frames are read
 - use AudioFileWritePackets() to write
 - advance the write position

Ch. 7 Audio Units: Generators, Effects, and Rendering

Where the Magic Happens

- Audio Units is lowest level of audio functionality that application developers need
- enables functionality that is impractical or impossible with the higher-level APIs
- increases difficulty of use and potential for serious error

How Audio Units Work

- works on streams of audio
- set up units, create connections, tell them to start processing
- analogous to the setup and flow of audio hardware
- types of audio units
 - generator units
 - create a stream of audio from some source, such as files, the network, or memory
 - instrument units
 - similar to generator units, produce a stream of synthesized audio from MIDI data
 - mixer units
 - combine multiple streams into one or more streams. the mix can be performed in two dimensions (for stereo panning) or in a simulated three-dimensional sound field
 - effect units
 - perform some sort of digital signal processing on a stream, usually producing an audible effect such as a reverb, a pitch change, noise filtering, and so on
 - converter units
 - perform transformations that are generally not meant to deliver user-audible effects
 - includes units to convert between different flavors of pcm (to change sample rate or bit depth, for example), adjust playback speed, and so on
 - output units
 - interface with audio input and/ or output hardware, enabling you to capture audio from input devices and play it out to output devices
 - name is misleading because these are potentially input/output units
- works with a pull model - software objects that need data pull samples from other objects using callbacks

Sizing Up the Audio Units

- type, subtype, and manufacturer are in AUComponent.h

Audio Unit Subtypes for Generator Units (Type `kAudioUnitType_Generator`)

- `kAudioUnitSubType_ScheduledSoundPlayer`
 - schedules audio to be played at a specific time
- `kAudioUnitSubType_AudioFilePlayer`
 - plays audio from a file
- `kAudioUnitSubType_NetReceive`
 - receives network audio from a corresponding `kAudioUnitSubType_NetSend` unit on another host or in another application

Audio Unit Subtypes for Instrument Units (Type `kAudioUnitType_MusicDevice`)

- `kAudioUnitSubType_DLSSynth`
 - multi-timbral music synthesizer that accepts MIDI commands
 - works with DLS or SoundFont formats

Audio Unit Subtypes for Mixer Units (Type `kAudioUnitType_Mixer`)

- `kAudioUnitSubType_MultiChannelMixer`
 - mixes any number of single or multi-channel input buses to one output bus
- `kAudioUnitSubType_StereoMixer`
 - mixes any number of mono or stereo input buses to one stereo output
- `kAudioUnitSubType_3DMixer`
 - mixes any number of mono or stereo input buses
 - mono inputs can be panned with 3D parameters
 - output is one bus of 2 to 8 channels
- `kAudioUnitSubType_MatrixMixer`
 - mixes any number of input and output buses, with any number of channels per bus
 - supports highly configurable mapping of inputs and outputs

Audio Unit Subtypes for Panner Units (Type `kAudioUnitType_Panner`)

- `kAudioUnitSubType_SphericalHeadPanner`
 - uses "spherical head" model to produce stereo output
- `kAudioUnitSubType_VectorPanner`
 - uses pan between adjacent channels in 3D space to create surround output
- `kAudioUnitSubType_SoundFilePanner`
 - uses "sound field" model to produce stereo output
- `kAudioUnitSubType_HRTFPanner`
 - uses head-related transfer function to produce stereo output

Audio Unit Subtypes for Effect Units (Type `kAudioUnitType_Effect`)

- `kAudioUnitSubType_Delay`
 - digital delay effect
- `kAudioUnitSubType_LowPassFilter`
- `kAudioUnitSubType_HighPassFilter`
- `kAudioUnitSubType_BandPassFilter`
- `kAudioUnitSubType_HighShelfFilter`
- `kAudioUnitSubType_LowShelfFilter`
- `kAudioUnitSubType_ParametricEQ`
- `kAudioUnitSubType_GraphicEQ`
- `kAudioUnitSubType_PeakLimiter`

- kAudioUnitSubType_DynamicsProcessor
- kAudioUnitSubType_MultiBandCompressor
- kAudioUnitSubType_MatrixReverb
- kAudioUnitSubType_SampleDelay
- kAudioUnitSubType_Pitch
- kAudioUnitSubType_AUFilter
 - adjusts gain for five bands of frequencies
- kAudioUnitSubType_NetSend
 - sends audio across network or between applications
- kAudioUnitSubType_Distortion
- kAudioUnitSubType_RogerBeep
 - produces a beep, similar to sound of walkie-talkie release, when the input level drops below a certain threshold for a certain amount of time

Audio Unit Subtypes for Converter Units (Type kAudioUnitType_FormatConverter)

- kAudioUnitSubType_AUConverter
 - uses audio converter services to perform LPCM conversions
- kAudioUnitSubType_Varispeed
 - changes playback speed, pitch-shifts audio as a result
 - faster = higher pitch
- kAudioUnitSubType_TimePitch
 - similar to varispeed
 - allows for independent control of playback speed and pitch
 - can play faster without changing pitch
- kAudioUnitSubType_DeferredRenderer
 - pulls input from a thread other than the caller
- kAudioUnitSubType_Splitter
 - splits one input bus into two identical output buses
- kAudioUnitSubType_Merger
 - merges two input buses into one output bus

Audio Unit Subtypes for Output Units (Type kAudioUnitType_Output)

- kAudioUnitSubType_GenericOutput
 - output unit not tied to audio hardware
 - can be used to perform software rendering of audio in connected units
- kAudioUnitSubType_SystemOutput
 - output to the device used for alerts and other UI sounds
- kAudioUnitSubType_DefaultOutput
 - output to the device selected in System Preferences -> Sound

- `kAudioUnitSubType_HALOutput`
 - input from and output to any supported audio device