

Memory Management

DSA C++

- general ideas behind basic malloc, realloc, free/new & delete
- different levels of memory management (hardware MMU, run time, garbage collection)
- external and internal fragmentation management

Sequential-fit

- for efficiency, use doubly linked list inside a block of memory, also maintaining a status field and a size field
- approaches:
 - first-fit: scan through the list and allocate the first block of memory large enough to meet the request
 - best-fit: scans through to find the block that is closest in fit to the requested size
 - worst-fit: use largest block in list so that the remaining portion is large enough to be used in future requests
 - next-fit: use the next available block that is sufficiently large
- analysis:
 - first-fit: most efficient
 - next-fit: comparable in speed but causes more external fragmentation
 - best-fit: causes the most fragmentation
 - worst-fit: avoids fragmentation in some ways
- fragmentation may be more dependent on the implementation than the method

Nonsequential-fit

- sequential-fit methods may become inefficient for large memory
- approach:
 - divide memory into an arbitrary number of lists, each holding blocks of the same size
 - larger blocks are split and the remainder is added to other lists or new lists are created if necessary
 - list sizes can be ordered into trees (Cartesian trees)
 - variants of first-fit and best-fit can be applied (i.e leftmost-fit and better-fit)
- approach:
 - programs have a limited number of sizes requested
 - use the above list approach with the most popular block sizes

- known as the adaptive exact-fit approach
- need to handle memory fragmentation:
 - occasional compaction
 - can also liquidate size list and rebuild

First-fit method (from TAOC 1)

- creates a linked list with the head pointing to the first available block
- each node has two fields for the related address
 - size of the memory contained at address P (i.e. number of words)
 - a link to the address of the next available block
- the last pointer is null
- algorithm:
 - overview: searches for an reserves the next available block or reports failure
 - variables: available (first available block), p, q
 - assumed preconditions:

```
link(p) returns a pointer to the next available block
available == link(location(available))
```

1. initialize:

```
q = location(available)
```

2. check for end of list:

```
p = link(q)
if (p == nullptr) { return or throw; }
```

3. check for adequate size:

```
if (size(p) >= requested_size) { continue; }
else { q = p; goto 2; }
```

4. reserve the size (and return the location -> edit)

```
k = size(p) - requested_size;
if (k == 0) { link(q) = link(p); return p; }
else { size(p) = k; return p + k; }
```

- NOTE: the algorithm terminates successfully on 4, reserving an area of length n beginning at location $p + k$
- improvement in performance: use the next fit method and rearrange links to maintain increasing order

Liberation with sorted list (from TAOC 1)

- follows from the First-fit method above but adds the assumption that if $\text{link}(p) \neq \text{nullptr}$ then $\text{link}(p) > p$ (i.e. the available list is sorted by memory location)
- adds a block of n consecutive cells beginning at p_0 to the available list with the assumption that none of the n cells is already available
- algorithm:
 - initialize

```
q = location(available)
```

2. advance p

```
p = link(q)
if (p == nullptr || p > p0) { goto 3; }
else { q = p; goto 2; }
```

3. check upper bound

```
if (p0 + n == p && p != nullptr) { n = n + size(p); link(p0) = link(p); }
else { link(p0) = p; }
```

4. check lower bound

```
// assumptions
size(location(available)) = 0;
if (q == location(available)) { throw; }

// otherwise
if (q + size(q) == p0) { size(q) = size(q) + n; link(q) = link(p0); }
else { link(q) = p0; size(p0) = n; }
```

Liberation with boundary tags (from TAOC 1)

- layout:

first word	second word	size - 2 words	last word
use flag size		block ...	

- assumptions:

- memory layout above, doubly linked (possibly in second and last word)
- available is the start of the doubly linked list
- places blocks of locations starting with p0 into the available list
- for the pool storage, locations m0 to m1 inclusive, assumes (tag meaning use flag, true meaning in use or reserved)

```
tag(m0 - 1) == tag(m1 + 1) == true;
```

- algorithm:

- check lower bound

```
if (tag(0 - 1)) { goto 3; }
```

2. delete lower area

```
p = p0 - size(p0 - 1);
p1 = link(p);
p2 = link(p + 1);
link(p1 + 1) = p2;
link(p2) = p1;
size(p) = size(p) + size(p0);
p0 = p;
```

3. check upper bound

```
p = p0 + size(p0);
if (tag(p)) { goto 5; }
```

4. delete upper area

```
p1 = link(p);
p2 = link(p + 1);
link(p1 + 1) = p2;
link(p2) = p1;
size(p0) = size(p0) + size(p);
p = p + size(p);
```

5. add to `available` list

```
size(p - 1) = size(p0);
link(p0) = available;
link(p0 + 1); = location(available);
link(available + 1) = p0;
available = p0;
tag(p0) = tag(p - 1) = false;
```

Basic memory allocation (from Modern Compiler Design)

- blocks: memory fragments released to user
- chunks: memory fragments controlled by allocator
- heap carves memory into chunks with the last byte of one chunk preceding the first byte of the following chunk

```

| chunk                                     |
|                                         |
|                                         | block
|                                         |
-----
| size   | use flag |          block          | size | use flag |   block
-----
size indicates size of chunk

```

- additional considerations - align block start on optimal alignment boundaries for a given data type -> this requires padding the size and flag segment or setting them back to ensure the block start location
- generalized algorithm

```
void* first_chunk = get_start_of_avail_mem();
void* past_end = first_chunk + size_of_avail_mem();

first_chunk->size = size_of_avail_mem();
first_chunk->free = true;

void* malloc(size_t size) {
    void* p = ptr_to_free_block_of_size(size);
    if (p != nullptr) { return p; }

    coalesce_free_chunks();
    void* p = ptr_to_free_block_of_size(size);
    if (p != nullptr) { return p; }

    // throws if fails
    solve_out_of_mem_cond(size);
}
```

```

        return malloc(size); // retry
    }

void free(void* p) {
    void *chunk = p - (sizeof(size_t) + sizeof(bool));
    chunk->free = true;
}

```

- compiler implementations:
 - linked lists
 - extensible arrays

Buddy Systems

- little external fragmentation and allows for compaction with little overhead
- generalized as a memory pool
- many different approaches
- can be combined with slab allocation to reduce fragmentation

Description (from TAOC 1)

- keep separate lists of available blocks of each size 2^k for $0 \leq k \leq m$
- entire pool consists of 2^m words with addresses 0 through $2^m - 1$
- starts with entire pool available
- when a block of memory is requested, if nothing of that size is available, a larger block is split in two to fulfill the request
- a split block forms two "buddies" and are rejoined when possible
- defined by fact that if the address and size of the first block is known, then the address and size of the second block is also known
- given the address of a block, x , with size 2^k

$$\text{buddy sub } k(x) = \{ x + 2^k, \text{ for } x \bmod 2^{(k+1)} == 0; \\ x - 2^k, \text{ for } x \bmod 2^{(k+1)} == 2^k; \}$$

- makes use of a one bit (or boolean) tag field in each block
- algorithm (reservation):
 - i. find block

```

j = min of range(k, m) if available_f[j] != location(available[j]); // inclusive r
// assumes available is sized  $2^j$  and is not empty
// throw if no such j exists

```

2. remove from list

```

l = available_b[j];
p = link_b(l);
available_b[j] = p;
link_f(p) = location(available[j]);
tag(l) = 0;

```

3. split required if necessary

```

if (j == k) { return address l; }

```

4. split

```

--j;
p = l + 2^j;
tag(p) = 1;
kval(p) = j;
link_f(p) = link_b(p) = location(available[j]);
available_f[j] = available_b[j] = p;
// splits a large block and registers the unused half in the available[j] list whi
goto 3;

```

- algorithm (liberation):
 - i. check if buddy is available

```

p = buddy sub k(l);
if (k == m || tag(p) == 0 || (tag(p) == 1 and kval(p) != k) { goto 3; }

```

2. combine with buddy

```

link_f(link_b(p)) = link_f(p);
link_b(link_f(p)) = link_b(p);
// removes block p from available[k] list
k = k + 1;
if (p < l) { l = p; goto 1; }

```

3. add l to `available` list

```

tag(l) = 1;
p = available_f[k];
link_f(l) = p;
link_b(p) = l;
kval(l) = k;

```

```
link_b(l) = location(available[k])
available_f[k] = l
// adds l to the available[k] list
```

Comparison of methods

- if first-fit and liberation with sorted list are used continuously, when the system tends to equilibrium with n reserved blocks on average, each equally likely to be the next freed block, and k in first-fit assuming non-zero values with probability p , then avg. number of available blocks tends to $\sim 1/2p \cdot n$
- using MIX/MMIX both the boundary tag system and the buddy system are fast for reservation and liberation, but the buddy system is slightly faster
- for known block size distributions, with equally likelihood that any block will be the next freed, the "distributed-fit method" results in much better memory utilization
 - partition memory into $n + n^{1/2} \cdot \lg(n)$
 - n = desired max number of blocks
 - each slot has a fixed size, different slots may have a different size, and any given slot has fixed boundaries and will be empty or contain a single allocated block

Caching (from wikipedia https://en.wikipedia.org/wiki/Cache_replacement_policies)

Belady's/clairvoyant algorithm

- discard information that will not be needed for the longest time in the future

FIFO

- removes element from the cache that has been in the cache for the longest time
- can use deque or queue

LIFO

- removes element from the cache that has been in the cache for the shortest time
- can use vector, stack, or deque

LRU (Least recently used)

- family of caching algorithms
 - LRU-K
 - ARC
 - MQ
 - 2Q
 - LRU-2
 - LRFU
 - LIRS
- involves keeping an age/access bit which is incremented on addition to the cache and incremented on cache access
- when a cache miss occurs, the cache item with the lowest age/access count is evicted and the element causing the cache miss is added

MRU (Most recently used)

- like LRU but discards the item with the highest age/access bit
- useful for repeated looping over an input in a reference pattern, random access patterns, and repeated scans over large data sets (aka cyclic access patterns)
- i.e. most useful when older items are more likely to be accessed

LFU (Least-frequently used)

- much like LRU but only tracks usage and removes the least used on a cache miss

Additional

- Pseudo-LRU (PLRU)
- Random replacement (RR)
- Segmented LRU (SLRU)
- LFU with Dynamic Aging (LFUDA)
- Low Inter-reference Recency Set (LIRS)
- Adaptive Replacement Cache (ARC)
- Clock with Adaptive Replacement (CAR)
- Multi Queue (MQ) caching algorithm | Multi Queue (MQ)
- Pannier: Container-based caching algorithm for compound objects

Garbage collection (wiki and compiler)

- reference counting: uses a usage count to determine when chunks are no longer needed -> requires concept of strong and weak references to avoid retain cycles
- mark and scan (mark and sweep): traverses chunk graphs and identifies reachable chunks, removing unreachable chunks
- two-space copying: maintains a from-space and a to-space, moving reachable chunks to the to-space (efficient but wastes 1/2 memory)
- higher level categorization:
 - tracing: what is commonly referred to as garbage collection. Generally, but not limited to, determining reachable objects and removing unreachable objects
 - reference counting: see above
 - escape analysis: method used to convert heap allocations to stack allocations to reduce garbage collector work -> generally done at compile time
- three varieties:
 - one-shot: runs through to completion
 - on-the-fly/incremental: run on each request for memory, mostly when more memory is needed or improvements to memory fragmentation/chunk structure is needed
 - concurrent: garbage collector runs on a separate processor

Mark and sweep

- naive:
 - each object has a use flag (always cleared except during collection)
 - traverses entire object tree (graph) from root and marks each pointed to object as in use
 - scans all memory from start to finish and frees any objects not marked as in-use
 - requires an entire scan of memory (some twice) which is a problem in paged memory
- tri-coloring:
 - white: candidate for recycling
 - black: no outgoing references but reachable from roots (not candidates)
 - gray: all objects reachable from root but yet to be scanned
 - procedure:
 - move object from gray set to black set
 - move each white object referenced by gray object to gray set
 - repeat last until gray set is empty
 - remove all remaining white objects

- has the advantage of being able to be performed on-the-fly or incrementally
- moving vs non-moving: may use two areas of memory and move reachable objects to separate area (aka compacting vs. non-compacting)
- copying vs. mark-and-sweep vs. mark-and-don't-sweep
- generational
- stop-the-world vs. incremental vs. concurrent
- precise vs. conservative and internal pointers

Cheney's algorithm

- stop and copy method of tracing garbage collection with a heap divided into two halves
- reclaims by:
 - if the object has not yet been moved to the to-space creates copy in to-space, replaces from-space version with a forwarding pointer to the to-space copy. Updates the object reference to refer to the new version in to-space.
 - if the object has already been moved to the to-space, simply update the reference from the forwarding pointer in from-space.

Generational garbage collector

- see above under high level categorization and tracing

Mark-compact algorithm

- uses mark-and-sweep and Cheney's copying algorithm to move reachable objects to another memory location (i.e. the beginning of the heap)
- used by GHC

Semi-space collector

- copying method mentioned above using from-space and to-space
- performs collection when to-space does not have adequate space to fulfill a request
- allocation costs are extremely low
- requires twice as much memory
- NOTE: Knuth and Compiler Techniques have more discussions in this area
- memory management
 - sequential-fit methods

- nonsequential-fit methods
- garbage collection
 - mark-and-sweep
 - copying methods
 - incremental
 - generational
- caching