

# A brief introduction to C++'s model for type- and resource-safety

---

- fundamental recurring errors:
  - resource leaks
  - access through an invalid pointer
  - memory corruption (dynamic memory is often reused and if accessed through an invalidated pointer it will be corrupted)
  - confusion between static and dynamic objects (free on static object, no free on dynamic object)
  - type-incorrect deallocation
- these problems have persisted for more than forty years starting in C
- garbage collection and finalizers are not an adequate solution

## Design constraints

---

- ideals:
  - i. Perfect static type safety
  - ii. Automatic release of general resources (perfect resource safety)
  - iii. No run-time overhead compared to good hand-crafted code (the zero-overhead principle)
  - iv. No restriction of the application domain compared to C and previous versions of C++
  - v. Compatibility with previous versions of C++ (long-term stability is a feature)

## Memory safety

---

- dangling pointers are a major problem
- ctor/dtor handles the resource management portion of the problem
- pointers can escape scopes through:
  - global variables
  - smart pointers
  - containers of pointers
  - objects holding pointers
  - more

- memory safety means all objects are allocated and deallocated once and only once and no object (pointer/reference etc) is accessed outside of this lifetime
- type safety is related to memory safety -> objects cannot be accessed through dangling pointers, no range errors, no null pointer access, etc
- use `owner` objects to manage pointers
  - includes pointer containers, smart pointers, etc
- dynamic ownership model:
  - attach an ownership bit to every pointer
  - set the ownership bit when a value returned from `new` is assigned
  - use `delete` whenever a pointer with the ownership bit is set goes out of scope
  - use `delete` before overwriting a pointer with an ownership bit set
  - clear the ownership bit when a pointer not returned from `new` is assigned
  - set the ownership bit when a pointer with the ownership bit set is assigned (and clear the ownership bit on the source; an object cannot have two owners)
- does not work in C++ (violates zero-overhead principal)
- static ownership model:
  - Mark every `T*` that is an owner as `owner<T*>` in the source code.
  - Let `new` return an `owner<T*>`
  - Make sure that every `owner<T*>` is deleted or transferred to another `owner<T*>`
  - Never assign a plain `T*` to an `owner<T*>`
- simply:

```
template<typename X> using owner = X;
```

- mainly used to assist programmers and static analysis tools
- see the GSL for more
- limitations include mixture with containers, returning from functions
- pointer safety:
  - Don't transfer a pointer to a local to where it could be accessed by a caller
  - A pointer passed as an argument can be passed back as a result
  - A pointer obtained from `new` can be passed back from a function as a result
- also apply to references and objects containing pointers or references

- pointer invalidation:
  - happens after delete/free
  - may happen in containers like vector after push\_back, clear, etc
- a dynamic invalidation model is not possible due to compatibility
- requires a static model and static checking
- be conservative when considering whether a function invalidates a pointer:
  - assume const functions do not invalidate
  - assume non-const functions might invalidate
  - can use `[[lifetime(const)]]` annotation
- limit smart pointer usage in interfaces to locations where pointer ownership is actually manipulated (for generality and performance)

## Resource Safety

---

- always delete new'ed objects only once, always consider where an exception may require
- use smart pointers to ensure safety
- polymorphism requires some form of pointer
- eliminate the use of pointers in containers because run-time polymorphism should not be used with them
- use move semantics to speed access in and out of these containers (enables transferring potentially large objects from one scope to another by "stealing" a handle, etc)

### ###!!!!IMPORTANT

- move semantics enables complete encapsulation of the management of non-scoped memory
- always place new's and deletes inside abstractions - no naked new's & no naked deletes
- RAII requires all failable resource acquisition to happen within a resource manager that can handle error reporting/management
- do not allow an owning pointer to be the only handle to an object in a context/scope where an exception can be thrown

## Leak-Freedom in C++ By Default

---

- lifetime guarantees:
  - no dangling/invalid dereferences (use after delete)

- no null pointer dereferences
- no leaks (always delete objects once and only once when they are no longer used)
- strategy:
  - i. prefer scoped lifetime by default (locals/members directly owned) - 80% of objects
  - ii. prefer `make_unique` & `unique_ptr` or a container if a an object must have its own lifetime and ownership can be unique (no ownership cycles) -> trees, lists - roughly 10% of objects
  - iii. use `make_shared` and `shared_ptr`s for an object with its own lifetime and shared ownership with owning cycles -> node-based DAGs (including trees that share references)
- don't use owning raw `*`'s means don't use explicit delete
- use `weak_ptr` to break cycles
- always prefer a non-static data member (non pointer) to model a "has-a" relationship
- always prefer a `unique_ptr` for a decoupled "has-a" relationship
  - separate/delayed/lazy/on-demand initialization
  - changeable component (polymorphism)
  - (manage move operations for non-null decoupled "has-a" members by writing your own)
- use a `const unique_ptr` for the pimpl idiom (compilation firewall)

```
template<class T>
using Pimpl = const unique_ptr<T>;

class MyClass {
    class Impl; // defined in .cpp
    Pimpl<Impl> pimpl;
    /*... Note: declare destructor and write it elsewhere ...*/
};
```

- use `const unique_ptr<Data[]>` for a dynamic array member (with `make_unique`)
- use `unique_ptr` for trees (both node children and root)
  - use a raw `Node*` for parents and enforce invariant `left->parent == this && right->parent == this`
  - releasing subtrees:
    - recursive:
      - automatic and correct
      - unbounded stack depth
    - iterative:
      - manual optimization
      - bounded stack depth
- use `unique_ptr` for doubly linked list forward links, root nodes and use raw `Node*` for reverse links

- used `shared_ptr` for children and root in Trees that hand out strong references to nodes use `shared_ptr` for each node's data
- use `shared_ptr` and `vector<shared_ptr>` for DAGs of heap objects (use `vector<Node*>` for parent references) use `shared_ptr` for each node's data
- for circular lists still use `unique_ptr`s but use a dummy node for the head and keep a reference in each node, when calling `next` either return `next.get()` or `head.get()`
- for cyclic graphs with nodes use `vector<Node*>` for children and parents in nodes and for roots, use `vector<unique_ptr>` nodes for unreachable node
- for factories
  - prefer `unique_ptr` + `make_unique` (if the object might not be shared by subsequent code)
  - use `shared_ptr` + `make_shared` if the object will definitely be shared
- for factories + caching:
  - use `shared_ptr`s with `weak_ptr`s
- avoiding issues (cycles):
  - don't pass an owner to unknown code
  - don't store an owner in a callback
- not all cycles can be broken with `weak_ptr`s
  - experimenting with `deferred_ptr`s and deferred containers to resolve naturally

## Lifetime Safety: Preventing Leaks and Dangling

---

- goal: eliminate leaks and dangling for `*/&/iterators/views/ranges`
  - no leaks or dangling
  - need efficiency of just using an address
  - cannot add run-time overhead
  - cannot add significant source code impact
  - cannot add requirements (for analysis) that adds excessive false positives
- uses `owner<>` type alias

### I. Approach and principles

---

- basic rules:
  - i. Prefer to allocate heap objects using (owning) `make_unique`/`make_shared` or containers

- ii. Otherwise, use `owner<T*>` for source/layout compatibility with old code. Each non-null `owner<>` must be deleted exactly once, or moved
- iii. Never dereference a null or invalid Pointer
- iv. Never allow an invalid Pointer to escape a function

## II. Informal overview and rationale

---

- shared owner (`shared_ptr`, `raw_shared_owner`) -> can't dangle
- unique owner (`containers`, `unique_ptr`, `owner`) -> can't dangle
- pointer non-owning (`raw *` and `&`, `iterators`, `ranges`, `views`) -> can dangle

### 1. Aliasing: taking addresses and dereferencing

- non-owning raw pointer
- local or member pointer is never an owner (only `owner<>` via `new`)
- `T&` is not `owner<T&>` and no conversion from `T*` to `owner<T*>`
- address of local variable is invalidated at end of scope
- address of member variable is invalidated at end of containing object's lifetime
- a pointer or object obtained by dereferencing has a lifetime dictated by the pointer/object it refers to

### 2. Invalidation by modifying Owners

- modifying ownership invalidates any pointer
- containers of containers

### 3. Branches

### 4. Loops

### 5. null

### 6. throw and catch

### 7. Calling functions (arguments and in/inout parameters)

### 8. Calling functions (return values and out/inout parameters)

### 9. Transferring ownership

### 10. Lifetime const

- NOTE: review each for all details -> contains applied examples at end

# A Mechanized Semantics for C++ Object Construction and Destruction, with Applications to Resource Management

---

- TODO: more here
- for now ->
- object construction:
  - i. Every object and each of its subobjects (if any) is constructed exactly once
  - ii. No object construction should rely on parts yet to be constructed, and no object destruction should rely on parts already destroyed
  - iii. Every object is destroyed in the exact reverse order it was constructed
- (means multi-part initialization is possible/fine)
- construction stages:
  - i. Unconstructed: Construction has not started yet.
  - ii. StartedConstructing: The construction of base-class subobjects has started, but not the fields.
  - iii. BasesConstructed: The base-class subobjects are completely constructed. Now constructing the fields or executing the constructor body.
  - iv. Constructed: The constructor body has returned, and destruction has not started yet.
  - v. StartedDestructing: The body of the destructor is executing or the fields are undergoing destruction.
  - vi. DestructingBases: The fields have been completely destructed. Bases are undergoing destruction.
  - vii. Destructed: All bases and fields have been destructed.
- properties:
  - run time invariants
  - progress
  - safety of field accesses and virtual function calls
  - evolution of construction states
  - object lifetimes
  - RAI
  - generalized dynamic types

# Writing Good C++14 By Default

---

- target guarantee: C++ code compiled in the safe subset is never the root cause of type/memory safety errors, except where explicitly annotated as unsafe

## Type safety overview

---

- GSL
  - byte (not char)
  - variant (tagged union)
- rules:
  - i. Don't use `reinterpret_cast`.
  - ii. Don't use `static_cast` downcasts. Use `dynamic_cast` instead.
  - iii. Don't use `const_cast` to cast away `const` (i.e., at all).
  - iv. Don't use C-style (T)expression casts that would perform a `reinterpret_cast`, `static_cast` downcast, or `const_cast`.
  - v. Don't use a local variable before it has been initialized.
  - vi. Always initialize a member variable.
  - vii. Avoid accessing members of raw unions. Prefer `variant` instead.
  - viii. Avoid reading from `varargs` or passing `vararg` arguments. Prefer variadic template parameters instead.

## Bounds safety overview

---

- GSL
  - `array_view` -> replaces passing an array + length
  - `string_view` -> convenience for char array
- rules:
  - i. Don't use pointer arithmetic. Use `array_view` instead.
  - ii. Only index into arrays using constant expressions.
  - iii. Don't use array-to-pointer decay.
  - iv. Don't use `std::` functions and types that are not bounds-checked.

## Lifetime safety overview

---

- use smart pointers (but don't overuse them)
- use raw pointers and references for efficiency, especially on the stack (locals, params, return values)



- GSL
  - Owner (can't dangle): owner<>, containers, smart pointers
  - Pointer (can dangle): \*, &, iterators, array\_view/string\_view, ranges
  - not\_null: wraps indirection and enforces non-null
- rules:
  - i. Prefer to allocate heap objects using make\_unique/make\_shared or containers.
  - ii. Otherwise, use owner<> for source/layout compatibility with old code. Each non-null owner<> must be deleted exactly once, or moved.
  - iii. Never dereference a null or invalid Pointer.
  - iv. Never allow an invalid Pointer to escape a function.

## Function bodies

- pointer to local -> do not access after scope (increase or decrease lifetime to match)
- address-of and pointer to pointer -> all fine but observe ownership rules
- dereferencing -> fine but observe ownership rules
- pointer from owner -> invalidated after owner modification
- be aware of temporaries for smart pointers and pointers from owners and dereferencing
- branches -> add the possibility of "or", be aware of invalidation in each branch
- loops -> like branches
- nullptr -> do not access if there is the possibility of null or not (consider not\_null above)
- try/catch -> treat catch as if it can be entered from any point in the try block
  - check invalidations in try block
  - any and all revalidations in the try block may not have been executed

## Function parameters

- for function definitions -> assume any pointer is valid
- in callers -> enforce that no invalid arguments are passed
- do not overuse smart pointers -> use for ownership modification

## Function return and in/out values

- minimize need to annotate ownership:
  - output pointers derived from input owner and pointer -> no annotation
  - no inputs = static pointer output -> no annotation (consider && refs)
  - anything else -> annotate
- a returned pointer is assumed to come from owner/pointer inputs

# Writing Good C++14

---

## Core Rules

---

- no leaks
  - scope every object in container/abstraction
  - RAI -> no naked new/delete
- no dangling pointers
  - distinguish owners
  - assume raw pointers are non-owners
  - catch all attempts for a pointer to escape its owner's scope
    - return, throw, out parameters, long-lived containers
  - something that holds an owner is an owner (i.e. containers of owners)
  - applies equally to references, containers of pointers, smart pointers
  - use owner $\triangleleft$  for low-level code and static analysis
  - rules:
    - Don't transfer to pointer to a local to where it could be accessed by a caller
    - A pointer passed as an argument can be passed back as a result
    - A pointer obtained from new can be passed back as a result as an owner
  - do not resort to just testing everywhere
- no type violations through pointers
  - eliminate range errors (array\_view, string\_view)
  - eliminate nullptr dereference
  - (casts)
  - (unions)

## Misuses of smart pointers

---

- used to:
  - represent ownership
  - avoid dangling pointers (unnecessary if ownership rules are followed)
- issues:
  - adds overhead (shared\_ptr)

- complicates function interfaces
- may not need pointers
  - unnecessary for scoped objects
  - necessary for:
    - OO interfaces -> polymorphism
    - need to change the object referred to -> ownership and polymorphism
    - (limit cost of passing large objects i.e. use of references)
- uses:
  - void f(T\*) -> use; no ownership transfer or sharing
  - void f(unique\_ptr) -> transfer unique ownership and use
  - void f(shared\_ptr) -> share ownership and use
- do not use a smart pointer in an interface if the function just uses the smart pointer